

PROBLEM SET 6

16822 GEOMETRY-BASED METHODS IN VISION (FALL 2022)

<https://piazza.com/cmu/fall2022/16822>

OUT: Nov. 16, 2022

DUE: Nov. 24, 2022 11:59 PM

Instructor: Shubham Tulsiani

TAs: Mosam Dabhi, Kangle Deng, Jenny Nan

START HERE: Instructions

- **Collaboration policy:** All are encouraged to work together BUT you must do your own work (code and write up). If you work with someone, please include their name in your write up and cite any code that has been discussed. If we find highly identical write-ups or code without proper accreditation of collaborators, we will take action according to university policies, i.e. you will likely fail the course. See the [Academic Integrity Section](#) detailed in the initial lecture for more information.
- **Late Submission Policy:** There are **no** late days for Problem Set submissions.
- **Submitting your work:**
 - We will be using Gradescope (<https://gradescope.com/>) to submit the Problem Sets. Please use the provided template. Submissions can be written in LaTeX. Regrade requests can be made, however this gives the TA the opportunity to regrade your entire paper, meaning if additional mistakes are found then points will be deducted. Each derivation/proof should be completed on a separate page. For short answer questions you **should** include your work in your solution.
- **Materials:** The data that you will need in order to complete this assignment is posted along with the writeup and template on Piazza.

For multiple choice or select all that apply questions, replace `\choice` with `\CorrectChoice` to obtain a shaded box/circle, and don't change anything else.

Instructions for Specific Problem Types

For “Select One” questions, please fill in the appropriate bubble completely:

Select One: Who taught this course?

- ☒ Shubham Tulsiani
- ☐ Deepak Pathak
- ☐ Fernando De la Torre
- ☐ Deva Ramanan

For “Select all that apply” questions, please fill in all appropriate squares completely:

Select all that apply: Which are scientists?

- ☒ Stephen Hawking
- ☒ Albert Einstein
- ☒ Isaac Newton
- ☐ None of the above

For questions where you must fill in a blank, please make sure your final answer is fully included in the given space. You may cross out answers or parts of answers, but the final answer must still be within the given space.

Fill in the blank: What is the course number?

16-822

1. [3 pts] Are these statements true or false?

- (a) Given a quaternion q , $-q$ represents its inverse rotation.
- (b) Given quaternions q_1 and q_2 , the composition of these 2 rotations is $q_1 + q_2$.
- (c) Given a unit quaternion q such that $\|q\| = 1$, there is a unique rotation it corresponds to.

a) False : $q^{-1} = q^* / \|q\|$
 b) False : it's the product
 c) True : a unit quaternion can be converted to unique rotation
 (However, if the question also implies one to one correspondence from rotation to quaternion then the answer would be False $\rightarrow +q, -q$)

2. [2 pts] Given a quaternion $q = [\sqrt{3}/2, 1/2, 0, 0]$, convert it to the axis-angle representation.

$q = (\cos \theta/2, \sin \theta/2 \hat{\omega}) = (\sqrt{3}/2, (1/2, 0, 0))$
 $\therefore \theta = 2 * \cos^{-1}(\sqrt{3}/2) = \pi/3$
 $\hat{\omega} = [\frac{1/2}{\sin \theta/2}, \frac{0}{\sin \theta/2}, \frac{0}{\sin \theta/2}]^T = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$
 $\therefore \|a\| = \pi/3$
 $\frac{a}{\|a\|} = \hat{\omega} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$
 $\therefore a = \begin{bmatrix} \pi/3 \\ 0 \\ 0 \end{bmatrix}$ Axis-Angle representation

3. [2 pts] In Slide 32 of Lecture 18, we introduced a parameterization of rotation that $\mathbf{R} = [\mathbf{r}_1, \mathbf{r}_2, \mathbf{r}_3] = F(\mathbf{v}_1, \mathbf{v}_2)$. Specifically,

$$\mathbf{r}_1 = \text{normalize}(\mathbf{v}_1), \quad \mathbf{r}_2 = \text{normalize}(\mathbf{v}_2 - (\mathbf{v}_2 \cdot \mathbf{r}_1)\mathbf{r}_1), \quad \mathbf{r}_3 = \mathbf{r}_1 \times \mathbf{r}_2.$$

given $\mathbf{v}_1, \mathbf{v}_2 \in \mathbb{R}^3$. What's the output $\mathbf{R}' = F(\mathbf{v}'_1, \mathbf{v}'_2)$ if $\mathbf{v}'_1 = \mathbf{v}_1, \mathbf{v}'_2 = -\mathbf{v}_2 + \mathbf{v}_1$? Represent the output $\mathbf{r}'_1, \mathbf{r}'_2, \mathbf{r}'_3$ using $\mathbf{r}_1, \mathbf{r}_2, \mathbf{r}_3$.

Given $\mathbf{v}'_1 = \mathbf{v}_1, \mathbf{v}'_2 = \mathbf{v}_1 - \mathbf{v}_2$
 $\therefore \mathbf{r}'_1 = \text{normalize}(\mathbf{v}'_1) = \text{normalize}(\mathbf{v}_1) = \mathbf{r}_1$ — ①
 $\mathbf{r}'_2 = \text{normalize}(\mathbf{v}'_2 - (\mathbf{v}'_2 \cdot \mathbf{r}'_1)\mathbf{r}'_1) = \text{normalize}(\mathbf{K})$
 where $\mathbf{K} = \mathbf{v}'_2 - (\mathbf{v}'_2 \cdot \mathbf{r}'_1)\mathbf{r}'_1$

$$\begin{aligned}
 \therefore K &= (v_1 - v_2) - ((v_1 - v_2)r_1)r_1 \\
 &= (v_1 - (v_1 \cdot r_1)r_1) - (v_2 - (v_2 \cdot r_1)r_1) \\
 &= (v_1 - v_1) - (v_2 - (v_2 \cdot r_1)r_1) \quad \because r_1 \text{ is just normalized } v_1
 \end{aligned}$$

$$\therefore K = -(v_2 - (v_2 \cdot r_1)r_1)$$

$$\therefore r_2' = \text{normalize}(K) = -\text{normalize}(v_2 - (v_2 \cdot r_1)r_1) = -r_2 \quad \text{--- (2)}$$

$$r_3' = r_1' \times r_2' = -r_1 \times r_2 = -r_3 \quad \text{--- (3)}$$

\therefore From (1), (2), (3)

$$r_1' = r_1, \quad r_2' = -r_2, \quad r_3' = -r_3$$

4. [2 pts] Assume two cameras on a stereo rig is looking at an object from 5 meters away. The baseline between the two cameras is 20 cm. We observe the disparity of a pixel on that object is 5 mm. How does the disparity change if

(a) the baseline between the two cameras is set to 40 cm?

(b) the camera is set to 10 meters away from the object but the baseline between the two cameras keeps the same?

a) disparity \propto baseline
 \therefore if baseline doubles, disparity doubles
 $\therefore d' = 2d = 10 \text{ mm}$

b) disparity $\propto \frac{1}{\text{distance}}$
 $\therefore d' = \frac{1}{2}d = 2.5 \text{ mm}$

5. [2 pts] How many potential keypoints does superpoint predict **at most** from a 128×128 image?

Superpoint involves getting $H/8 \times W/8$ cells of 8×8 pixels (+1 dustbin)
 The softmax probabilities are used & then reshaped back to $H \times W$ & top n keypoints are taken.
 $\therefore \text{Max keypoints} = 128 \times 128 = 16384$

6. [3 pts] Given corresponding 3D points $\{(\mathbf{p}_i, \mathbf{q}_i)\} \in (\mathbb{R}^3, \mathbb{R}^3)$, we want to find the optimal rigid transform. Assume the points are already mean-centered so we only need to find the rotation. For every pair of points $(\mathbf{p}_i, \mathbf{q}_i)$, both points are located on one of the x , y and z axes. Namely,

$$\begin{aligned} \forall i, \quad & \mathbf{p}_i = [u_i, 0, 0], \quad \mathbf{q}_i = [v_i, 0, 0] \\ \text{or} \quad & \mathbf{p}_i = [0, u_i, 0], \quad \mathbf{q}_i = [0, v_i, 0] \\ \text{or} \quad & \mathbf{p}_i = [0, 0, u_i], \quad \mathbf{q}_i = [0, 0, v_i] \end{aligned}$$

(a) [1 pts] What are the possible rotations intuitively?

(b) [2 pts] Prove they are the **only** answers mathematically.

a) Intuitively, rotation angle along each axes should be 0 or π
 $\therefore R = \begin{bmatrix} \pm 1 & 0 & 0 \\ 0 & \pm 1 & 0 \\ 0 & 0 & \pm 1 \end{bmatrix}$

b) Contd.

7. [6 pts] We introduced 4 learning-based reconstruction methods in the lecture: PixelNeRF, NeRFormer, SRT, GBT.

- (a) [2 pts] Which of them use volumetric rendering?
- (b) [2 pts] Which of them use projection-based features?
- (c) [2 pts] Which of them use learning-based aggregation?

Feel free to write your reasoning in case you are unsure.

- a) PixelNeRF, NeRFormer
- b) PixelNeRF, NeRFormer
- c) NeRFormer, SRT, GBT

$$b) \quad M = \mathbf{Q} \mathbf{P}^T = \begin{bmatrix} \sum_x u_i v_i & 0 & 0 \\ 0 & \sum_y u_i v_i & 0 \\ 0 & 0 & \sum_z u_i v_i \end{bmatrix} = \begin{bmatrix} a & 0 & 0 \\ 0 & b & 0 \\ 0 & 0 & c \end{bmatrix}$$

$$\text{SVD}(M) = \mathbf{U} \Sigma \mathbf{V}^T \\ = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} a & 0 & 0 \\ 0 & b & 0 \\ 0 & 0 & c \end{bmatrix} \begin{bmatrix} \text{sign}(a) & 0 & 0 \\ 0 & \text{sign}(b) & 0 \\ 0 & 0 & \text{sign}(c) \end{bmatrix}$$

$$\therefore R = \mathbf{U} \mathbf{V}^T = \begin{bmatrix} \pm 1 & 0 & 0 \\ 0 & \pm 1 & 0 \\ 0 & 0 & \pm 1 \end{bmatrix}$$

Collaboration Questions Please answer the following:

1. Did you receive any help whatsoever from anyone in solving this assignment?

☐ Yes

☒ No

- If you answered 'Yes', give full details:
- (e.g. "Jane Doe explained to me what is asked in Question 3.4")

2. Did you give any help whatsoever to anyone in solving this assignment?

☐ Yes

☒ No

- If you answered 'Yes', give full details:
- (e.g. "I pointed Joe Smith to section 2.3 since he didn't know how to proceed with Question 2")

3. Did you find or come across code that implements any part of this assignment ?

☐ Yes

☒ No

- If you answered 'Yes', give full details: No
- (book & page, URL & location within the page, etc.).